

Manutenção e Evolução de Software

Marco Túlio Valente

mtov@dcc.ufmg.br

DCC - UFMG

How We Refactor, and How We Know

Emerson Murphy-Hill, Chris Parnin, Andrew P. Black
ICSE 2009, Best Paper

Introduction

- Refactoring is the process of changing the structure of a program without changing the way that it behaves.
- To help put refactoring research on a sound scientific basis, we draw conclusions using four data sets spanning:
 - More than 13 000 developers
 - 240 000 tool-assisted refactorings
 - 2500 developer hours
 - 3400 version control commits.
- Using these data, we cast doubt on several previously stated assumptions about how programmers refactor, while validating others

The Data that We Analyzed (1/3)

- The work described in this paper is based on four sets of data
- The first set we will call **Users**
 - It was originally collected in the latter half of 2005 by Murphy and colleagues who used the Mylyn Monitor tool to capture and analyze fine-grained usage data from 41 volunteer programmers in the wild using Eclipse
 - These data capture an average of 66 hours of development time per programmer
 - About 95 percent of the programmers wrote in Java.
 - The data include information on which Eclipse commands were executed, and at what time

The Data that We Analyzed (2/3)

- The second set of data we will call **Everyone**
 - It is publicly available from the Eclipse Usage Collector
 - It includes data from every user of the Eclipse Ganymede release who consented to an automated request to send the data back to the Eclipse Foundation.
 - These data aggregate activity from over 13 000 Java developers between April 2008 and January 2009, but also include non-Java developers.
 - The data count how many programmers have used each Eclipse command, including refactoring commands, and how many times each command was executed.
- We know of no other research that has used these data for characterizing programmer behavior

The Data that We Analyzed (3/3)

- The third set of data we will call **Toolsmiths**
 - It includes refactoring histories from 4 developers who primarily maintain Eclipse's refactoring tools.
 - These data include detailed histories of which refactorings were executed, when they were performed, and with what configuration parameters
- The fourth set of data we will call **Eclipse CVS**
 - It is the version history of the Eclipse and JUnit extracted from CVS repositories.

Findings on Refactoring Behavior

Toolsmiths and Users Differ

| Refactoring Tool | Toolsmiths | | | | Users | | | | Everyone | |
|------------------------------|------------|-------|---------|-----------|-------|-------|---------|-----------|----------|-------|
| | Uses | Use % | Batched | Batched % | Uses | Use % | Batched | Batched % | Uses | Use % |
| Rename | 670 | 28.7% | 283 | 42.2% | 1862 | 61.5% | 1009 | 54.2% | 179871 | 74.8% |
| Extract Local Variable | 568 | 24.4% | 127 | 22.4% | 322 | 10.6% | 106 | 32.9% | 13523 | 5.6% |
| Inline | 349 | 15.0% | 132 | 37.8% | 137 | 4.5% | 52 | 38.0% | 4102 | 1.7% |
| Extract Method | 280 | 12.0% | 28 | 10.0% | 259 | 8.6% | 57 | 22.0% | 10581 | 4.4% |
| Move | 147 | 6.3% | 50 | 34.0% | 171 | 5.6% | 98 | 57.3% | 13208 | 5.5% |
| Change Method Signature | 93 | 4.0% | 26 | 28.0% | 55 | 1.8% | 20 | 36.4% | 4764 | 2.0% |
| Convert Local To Field | 92 | 3.9% | 12 | 13.0% | 27 | 0.9% | 10 | 37.0% | 1603 | 0.7% |
| Introduce Parameter | 41 | 1.8% | 20 | 48.8% | 16 | 0.5% | 11 | 68.8% | 416 | 0.2% |
| Extract Constant | 22 | 0.9% | 6 | 27.3% | 81 | 2.7% | 48 | 59.3% | 3363 | 1.4% |
| Convert Anonymous To Nested | 18 | 0.8% | 0 | 0.0% | 19 | 0.6% | 7 | 36.8% | 269 | 0.1% |
| Move Member Type to New File | 15 | 0.6% | 0 | 0.0% | 12 | 0.4% | 5 | 41.7% | 838 | 0.3% |
| Pull Up | 12 | 0.5% | 0 | 0.0% | 36 | 1.2% | 4 | 11.1% | 1134 | 0.5% |
| Encapsulate Field | 11 | 0.5% | 8 | 72.7% | 4 | 0.1% | 2 | 50.0% | 1739 | 0.7% |
| Extract Interface | 2 | 0.1% | 0 | 0.0% | 15 | 0.5% | 0 | 0.0% | 1612 | 0.7% |
| Generalize Declared Type | 2 | 0.1% | 0 | 0.0% | 4 | 0.1% | 2 | 50.0% | 173 | 0.1% |
| Push Down | 1 | 0.0% | 0 | 0.0% | 1 | 0.0% | 0 | 0.0% | 279 | 0.1% |
| Infer Generic Type Arguments | 0 | 0.0% | 0 | - | 3 | 0.1% | 0 | 0.0% | 703 | 0.3% |
| Use Supertype Where Possible | 0 | 0.0% | 0 | - | 2 | 0.1% | 0 | 0.0% | 143 | 0.1% |
| Introduce Factory | 0 | 0.0% | 0 | - | 1 | 0.0% | 0 | 0.0% | 121 | 0.1% |
| Extract Superclass | 7 | 0.3% | 0 | 0.0% | * | - | * | * | 558 | 0.2% |
| Extract Class | 1 | 0.0% | 0 | 0.0% | * | - | * | * | 983 | 0.4% |
| Introduce Parameter Object | 0 | 0.0% | 0 | - | * | - | * | * | 208 | 0.1% |
| Introduce Indirection | 0 | 0.0% | 0 | - | * | - | * | * | 145 | 0.1% |
| Total | 2331 | 100% | 692 | 29.7% | 3027 | 100% | 1431 | 47.3% | 240336 | 100% |

Table 1. Refactoring tool usage in Eclipse. Some tool logging began in the middle of the *Toolsmiths* data collection (shown in light grey) and after the *Users* data collection (denoted with a *).

Programmers Repeat Refactorings

- We hypothesize that when programmers perform a refactoring, they typically perform several refactorings of the same kind within a short time period.
- We used the Toolsmiths and the Users data to measure the temporal proximity of refactorings to one another.
- We say that refactorings of the same kind that execute within 60 seconds of each another form a batch.
- In Table 1, each “Batched” column indicates the number of refactorings that appeared as part of a batch
- In total, we see that 30% of Toolsmiths refactorings and 47% of Users refactorings appear as part of a batch.

Configuration Options

- Refactoring tools are typically of two kinds:
 - They either force the programmer to provide configuration information, such as whether a newly created method should be public or private
 - Or they quickly perform a refactoring without allowing any configuration
- “Change Frequency” refers to how often a user used a configuration option other than the default.
- The data suggest that refactoring tools are configured very little: the overall mean change frequency for these options is just under 10%.

Configuration Options

| Refactoring Tool | Configuration Option | Default Value | Change Frequency |
|------------------------|---|---------------|------------------|
| Extract Local Variable | Declare the local variable as 'final' | false | 5% |
| Extract Method | New method visibility | private | 6% |
| | Declare thrown runtime exceptions | false | 24% |
| | Generate method comment | false | 9% |
| Rename Type | Update references | true | 3% |
| | Update similarly named variables and methods | false | 24% |
| | Update textual occurrences in comments and strings | false | 15% |
| | Update fully qualified names in non-Java text files | true | 7% |
| Rename Method | Update references | true | 0% |
| | Keep original method as delegate to renamed method | false | 1% |
| Inline Method | Delete method declaration | true | 9% |

Table 2. Refactoring tool configuration in Eclipse from *Toolsmiths*.

Floss Refactoring is Common

- In previous work, we distinguished two tactics that programmers use when refactoring: floss refactoring and root-canal refactoring.
- During floss refactoring, the programmer uses refactoring as a means to reach a specific end, such as adding a feature or fixing a bug.
- Thus, during floss refactoring the programmer intersperses refactoring with other kinds of program changes to keep code healthy.
- Root-canal refactoring, in contrast, is used for correcting deteriorated code and involves a protracted process consisting of exclusive refactoring.

Floss Refactoring is Common

- For convenience, we let a programming session be the period of time between consecutive commits to CVS by a single programmer.
- In a particular session, if a programmer both refactors and makes a semantic change, then we say that that the programmer is floss refactoring.
- If a programmer refactors during a session but does not change the semantics of the program, then we say that the programmer is root-canal refactoring
- Commits indicating floss refactoring would account for 30% of commits while commits indicating root-canal would account for only 3% of commits

Floss Refactoring is Common

- In no more than 10 out of 2671 commits did programmers use a refactoring tool without also manually editing their program.
- In other words, in less than 0.4% of commits did we observe the possibility of root-canal refactoring using only refactoring tools.

Medium and Low-level Refactorings

- We divided the refactorings into three levels — High, Medium and Low.
- High level refactorings are those that change the signatures of classes, methods, and fields;
 - Examples: RENAME CLASS, MOVE STATIC FIELD, and ADD PARAMETER.
- Medium level refactorings are those that change the signatures of classes, methods, and fields and also significantly change blocks of code
 - Example: EXTRACT METHOD, INLINE CONSTANT, and CONVERT ANONYMOUS TYPE TO NESTED TYPE.

Medium and Low-level Refactorings

- Low level refactorings are those that make changes to only blocks of code
 - Example: EXTRACT LOCAL VARIABLE, RENAME LOCAL VARIABLE, and ADD ASSERTION

Medium and Low-level Refactorings




Figure 2. Refactorings over 40 sessions.

Refactorings are Frequent

- In the Toolsmiths data, we found that refactoring activity occurred throughout the Eclipse development cycle.
 - In 2006, an average of 30 refactorings took place each week;
 - In 2007, there were 46 refactorings per week.
 - Only two weeks in 2006 did not have any refactoring activity, and one of these was a winter holiday week.
 - In 2007, refactoring occurred every week.

Refactorings are Frequent

- In the Users data set, we found refactoring activity distributed throughout the programming sessions.
 - Overall, 41% of programming sessions contained refactoring activity.
 - More interestingly, sessions that did not have refactoring activity contained an order of magnitude fewer edits than sessions with refactoring, on average.
 - The sessions that contained refactoring also contained, on average, 71% of the total edits made by the programmer

Refactorings with and without tools

- Comparing these results, we inferred that the EXTRACT METHOD tool is underused: the refactoring is instead being performed manually.
- However, it is unclear what other refactoring tools are underused. Moreover, there may be some refactorings that must be performed manually because no tool yet exists.
- We suspect that the reason that some kinds of refactoring — especially RENAME— are more often performed with tools is because these tools have simpler user interfaces.

Discussion: Tool-Usage Behavior

- Several of our findings have reflected on the behavior of programmers using refactoring tools.
- For example, our finding about how toolsmiths differ from regular programmers in terms of refactoring tool usage (Section 3.1) suggests that most kinds of refactorings will not be used as frequently as the toolsmiths hoped, when compared to the frequently used RENAME refactoring.
- For the toolsmith, this means that improving the underused tools or their documentation (especially the tool for EXTRACT LOCAL VARIABLE) may increase tool use

Discussion: Tool-Usage Behavior

- Our findings about the batching of refactorings provides evidence that tools that force the programmer to repeatedly select, initiate, and configure can waste programmers' time.
- This was in fact one of the motivations for Murphy-Hill and Black's refactoring cues, a tool that allows the programmer to select several program elements for refactoring at one time

Discussion: Refactoring Practice

- Several of our findings bolster existing evidence about refactoring practice across a large population of programmers
- Second, our finding that floss refactoring is a more frequently practiced refactoring tactic than root-canal refactoring confirms that floss refactoring, in addition to being recommended by experts, is also popular among programmers
- For toolsmiths, this means that refactoring tools should support flossing by allowing the programmer to switch quickly between refactoring and other development activities, which is not always possible with existing refactoring tools, such as those that force the programmer's attention away from the task at hand with modal dialog boxes

Discussion: Refactoring Practice

- Toolsmiths need to explore alternative interfaces and identify common refactoring workflows, such as reminding users to `EXTRACT LOCAL VARIABLE` before an `EXTRACT METHOD` or finding a easy way to combine these refactorings: the goal should be to encourage and support programmers in taking full advantage of refactoring tools.

Conclusions and Contributions

- The RENAME refactoring tool is used much more frequently by ordinary programmers than by the developers of refactoring tools
- About 40% of refactorings performed using a tool occur in batches
- About 90% of configuration defaults of refactoring tools remain unchanged when programmers use the tools
- Messages written by programmers in commit logs do not reliably indicate the presence of refactoring
- Programmers frequently floss refactor, that is, they interleave refactoring with other types of programming activity

Conclusions and Contributions

- About half of refactorings are not high-level, so refactoring detection tools that look exclusively for high-level refactorings will not detect them
- Refactorings are performed frequently
- Almost 90% of refactorings are performed manually, without the help of tools
- The kind of refactoring performed with tools differs from the kind performed manually